

# ClickHouse – Open Source OLAP DBMS

introduction & cibersecurity applications

João Pinheiro - 2022

{ Overview }

- Open Source OLAP database, created by Yandex;
- True column-oriented approach (no extra data stored with values);
- Optimized data compression;
- Designed to work with regular hard drives;
- Distributed query support via sharding;
- Suitable for online queries – sub-second latencies;
- Data replication (async multi-master);
- Queriable via SQL Query language;

- Small to large datasets with billions or trillions of rows;
- Data is organized in “wide” tables, comprised of many columns;
- The result dataset of a given query is a limited set of rows;
- Common query operations are aggregation queries;
- Results must be obtained in seconds or less;

- ClickHouse is fast. And by fast, I mean FAST.
- Easy to deploy & maintain, including scaling horizontally and backup;
- Multiple ingestion sources (SQL, CSV/TSV, S3, Kafka, etc);
- Supports foreign tables (MySQL, MongoDB, PostgreSQL, etc);
- Rich data types, suitable for many different usage scenarios;
- Multiple table engines suitable for different applications;
- Easy to use (if you know SQL);
- Integrates with existing dashboarding tools like grafana;
- Did I mention how fast it is?

- No “proper” transactions;
- Datasets are mostly “insert-read”;
- Limited functionality on updates & deletes and they are processed as batch operations;
- Limited performance on point queries returning a single row by key;
- No table relations;
- Works better with denormalized data;
- Inserts should be batched – it takes roughly the same time to insert 10 rows or 1000 rows;

**Integer:** UInt{8,16,32,64,128,256}, Int{8,16,32,64,128,256}

**Variable precision:** Float32, Float64

**Fixed precision:** Decimal

**String/Char/BLOB:** String, FixedString(n)

**Date/Time:** Date, Date32, DateTime, DateTime64

**Geo:** Point, Ring, Polygon, MultiPolygon

**Misc:** Boolean, UUID, LowCardinality, Enum, Array, JSON\*, Nested

**Network:** Ipv4, IPv6

### **\*MergeTree engine**

- Quick inserts with background merging;
- Data replication with Replicated\* engines;
- Optional partitioning via partitioning key;
- Supports primary key/sorting key;
- Advanced features: sampling & TTL;

### **Memory engine**

- Uses RAM without compression;
- No overhead on read operations;
- Does not support indexes;
- Non-persistent;

### **\* Log engine**

- Small tables, up to 1M records;
- Supports HDFS and S3 as file systems;
- Does not support indexes nor mutations;

### **Integration engines**

- ODBC/JDBC;
- MySQL, MongoDB, PostgreSQL, RocksDB;
- HDFS, S3;
- Kafka, RabbitMQ, File;



**{ showcase }**

importing a dataset for quick analysis

# Ingesting data from a well-known leak: install ClickHouse

{ Showcase }

```
sudo apt-get install -y apt-transport-https ca-certificates dirmngr
sudo apt-key adv --keyserver hkp://keyserver.ubuntu.com:80 --recv
8919F6BD2B48D754

echo "deb https://packages.clickhouse.com/deb stable main" | sudo tee \
/etc/apt/sources.list.d/clickhouse.list
sudo apt-get update

sudo apt-get install -y clickhouse-server clickhouse-client

sudo service clickhouse-server start
clickhouse-client # or "clickhouse-client --password" if you've set up a
password.
```

## Machine:

Hetzner Cloud VPS  
Ubuntu 22.04  
3x AMD Epyc 2.5Ghz  
4GB RAM  
80GB SSD

## Source to be ingested:

**19GB** of CSV files



# Ingesting data from a well-known leak: create a table

{ Showcase }

```
CREATE TABLE IF NOT EXISTS tap(  
  MDM_ID String,  
  BIRTH_DATE String,  
  SALUTATION String,  
  COMMUNICATION_LANGUAGE LowCardinality(String),  
  FIRST_NAME String,  
  LAST_NAME String,  
  FULL_NAME String,  
  NATIONALITY LowCardinality(String),  
  GENDER LowCardinality(String),  
  ADDRESS_CITY LowCardinality(String),  
  ADDRESS_COUNTRY LowCardinality(String),  
  ADDRESS_DETAIL String,  
  ADDRESS_REGION String,  
  ADDRESS_ZIPCODE String,  
  
  (...)  
  
  STATUS LowCardinality(String),  
  
  INDEX idx01 (ADDRESS_CITY) TYPE set(0) GRANULARITY 1  
  
) ENGINE = MergeTree()  
  ORDER BY (ADDRESS_COUNTRY, FIRST_NAME, LAST_NAME);
```

## Machine:

Hetzner Cloud VPS  
Ubuntu 22.04  
3x AMD Epyc 2.5Ghz  
4GB RAM  
80GB SSD

## Source to be ingested:

**19GB** of CSV files



# Ingesting data from a well-known leak: ingest data

{ Showcase }

## Import bash script

```
for FILENAME in /data/src/Customer_*.csv; do
    echo $FILENAME
    clickhouse-client --query="INSERT INTO tap FORMAT CSV" \
        --format_csv_delimiter="|" < $FILENAME
done
```

```
root@ubuntu-4gb-fsn1-2:/data# time ./clickhouse-import.sh
```

```
(...)
```

```
real    12m8.895s
user    9m23.545s
sys     0m34.991s
```

```
root@ubuntu-4gb-fsn1-2:/data#
```

## Give it 2-3 Minutes to settle and check disk usage

```
root@ubuntu-4gb-fsn1-2:/data# du -h /var/lib/clickhouse/
```

```
(...)
```

```
3.9G    /var/lib/clickhouse/
```

```
root@ubuntu-4gb-fsn1-2:/data#
```

## Machine:

Hetzner Cloud VPS  
Ubuntu 22.04  
3x AMD Epyc 2.5Ghz  
4GB RAM  
80GB SSD

## Source to be ingested:

**19GB** of CSV files

## ClickHouse disk usage:

**3.9GB**



BLACKSHIELD

# Ingesting data from a well-known leak: query time!

{ Showcase }

## How many records

```
ubuntu-4gb-fsn1-2 :) select count(*) from tap;
```

```
SELECT count(*)  
FROM tap
```

```
Query id: 0f315a2e-2419-4254-b8d7-83a891480cd5
```

```
count()  
77629463
```

```
1 row in set. Elapsed: 0.003 sec.
```

77 Million records in  
12m8s:

~ 6.4M inserts/minute  
~10600 inserts/second

## How many unique e-mail addresses

```
ubuntu-4gb-fsn1-2 :) select count(distinct(EMAIL_DESCRIPTION)) from tap;
```

```
SELECT countDistinct(EMAIL_DESCRIPTION)  
FROM tap
```

```
Query id: 02f3174e-1e29-470d-b52e-7546ad952b41
```

```
uniqExact(EMAIL_DESCRIPTION)  
6075465
```

```
1 row in set. Elapsed: 4.170 sec. Processed 77.63 million rows, 2.48 GB  
(18.62 million rows/s., 593.82 MB/s.)
```

How many records for someone with 'Manuel' on the name and from Gouveia city (case sensitive)

```
ubuntu-4gb-fsn1-2 :) select count(*) from tap where ADDRESS_CITY='Gouveia' and FULL_NAME like '%Manuel %';
```

```
SELECT count(*)  
FROM tap  
WHERE (ADDRESS_CITY = 'Gouveia') AND (FULL_NAME LIKE '%Manuel %')
```

Query id: a49fb40d-b44f-4120-9ef2-9894420ccabe

count()
45

1 row in set. Elapsed: 0.417 sec. Processed 6.29 million rows, 144.64 MB (15.09 million rows/s., 346.86 MB/s.)

### Top10 of records per Country

```
ubuntu-4gb-fsn1-2 :) select ADDRESS_COUNTRY,count(*) as total from tap group by(ADDRESS_COUNTRY) order by total desc limit 10;
```

(...)

ADDRESS_COUNTRY	total
BR	21682463
PT	17130056
US	5911385
FR	4952701
GB	4057600
ES	3555442
DE	3146254
	2603944
IT	2578419
CH	1961251

10 rows in set. Elapsed: 0.210 sec. Processed 77.63 million rows, 81.77 MB (369.45 million rows/s., 389.14 MB/s.)

### Export all e-mail addresses to a CSV file

```
ubuntu-4gb-fsn1-2 :) select distinct(EMAIL_DESCRIPTION) from tap into outfile 'emails.csv' format CSV;
```

```
SELECT DISTINCT EMAIL_DESCRIPTION  
FROM tap  
INTO OUTFILE 'emails.csv'  
FORMAT CSV
```

```
Query id: 245fa26a-b302-4456-8570-194c52885974
```

```
6075465 rows in set. Elapsed: 5.542 sec. Processed 77.63 million rows, 2.48 GB (14.01 million rows/s., 446.80 MB/s.)
```

```
ubuntu-4gb-fsn1-2 :)
```



**{ showcase }**

querying network events

### Event table

```
CREATE TABLE IF NOT EXISTS Events
(
    Id UUID,
    Created DateTime32,
    Address IPv4,
    Domain String
) ENGINE = MergeTree()
ORDER BY (Id, Created)
PRIMARY KEY Id;
```

#### Machine:

Hetzner Cloud VPS  
Ubuntu 22.04  
3x AMD Epyc 2.5Ghz  
4GB RAM  
80GB SSD

#### Dataset Size:

149.997.000 Events

### Sample data

Id	Created	Address	Domain
911008ec-7681-46a2-8000-066509e1f838	2022-03-04 10:20:40	76.172.30.184	ggpewp.vhn
f4804030-56b2-4776-8000-1a67631308b2	2021-08-13 22:51:54	146.193.180.94	ykkxgckrtscs.vbr
1889327e-8f14-4b61-8000-20da5fef938e	2021-07-25 09:03:05	195.33.251.217	vxscmq.fim
7ff7fddf-4fce-47b3-8000-4a7330474c55	2021-02-11 18:42:10	115.174.75.31	kdkrgkgwna.eno
eeb954eb-8c2b-4187-8000-55b3c318b599	2022-08-06 12:40:44	217.63.148.2	zuulwdjppqkn.aqr
bb6ea2f1-cf55-46a5-8000-832dc58c428f	2021-01-13 23:46:01	7.226.57.213	mouboknjhukz.loy



List first 10 records whose IP is within the subnet 10.0.0/16

```
ubuntu-4gb-fsn1-2 :) SELECT Created, Domain, Address FROM Events  
WHERE isIPAddressInRange(IPv4NumToString(Address), '10.0.0.0/16')=1  
LIMIT 10;
```

(...)

Created	Domain	Address
2022-08-17 06:31:30	pcxjttjhym.acr	10.0.218.42

10 rows in set. Elapsed: 0.041 sec. Processed 712.70 thousand rows, 4.00 MB (17.25 million rows/s., 96.83 MB/s.)



Retrieve the first 10 records that match:

- Address matches both 10.0.0.0/8 and 10.0.10.0/16
- TLD of domain starts with letter 'a'
- Address has a 2 somewhere in it
- Created between 01-01-2022 and 01-10-2022
- Older events first

```
ubuntu-4gb-fsn1-2 :) SELECT Created, Domain, IPv4NumToString(Address) AS addr
                     FROM Events
                     WHERE isIPAddressInRange(addr, '10.0.0.0/8')=1
                     AND Domain like '%.a%'
                     AND addr like '%2%'
                     AND isIPAddressInRange(addr, '10.0.10.0/16')=1
                     AND Created BETWEEN '2022-01-01 00:00:00' AND '2022-10-01 00:00:00'
                     ORDER BY Created
                     LIMIT 10;
```



(...)

Created	Domain	addr
2022-01-02 15:06:59	mhzebj.anh	10.0.18.244
2022-02-20 02:24:42	apcquscimh bv.abu	10.0.48.239
2022-03-02 13:39:15	qzlc lgs ljk v.aqz	10.0.55.223
2022-03-19 05:04:59	crcijzzjbdui.alo	10.0.140.243
2022-03-27 20:09:38	fenyyesjiq.aol	10.0.241.240
2022-04-06 07:52:18	bsyfnrympr.amp	10.0.122.175
2022-04-06 21:30:37	dvskgbfnab.anq	10.0.207.63
2022-05-01 20:22:16	wplaxuv.arr	10.0.228.97
2022-05-05 18:32:27	bxmdouxwa.agd	10.0.15.212
2022-05-05 19:56:59	gqeql dkip.adg	10.0.2.241

10 rows in set. Elapsed: 7.897 sec. Processed 150.00 million rows, 4.50 GB (18.99 million rows/s., 569.80 MB/s.)

**{ thank you }**  
questions?

**Project page:** <https://clickhouse.com/>

**Tutorials and datasets:** <https://clickhouse.com/docs/en/getting-started/example-datasets/>

**Playground with datasets:** <https://clickhouse.com/docs/en/getting-started/playground>

**Altinity Blog:** <https://altinity.com/blog/>

**Curated list of resources:** <https://github.com/korchasa/awesome-clickhouse>

**Interactive Benchmark:** <https://benchmark.clickhouse.com/>